

Loading Libraries

```
In [43]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import statsmodels.api as sm
import seaborn as sns
sns.set()
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

import warnings
warnings.filterwarnings('ignore')
```

Load Data

```
In [2]: f_birth = pd.read_csv(r"E:\Data Science Projects\Project Files\Timeseries\Daily total female births in California, 1959 new.csv",
                             sep=",", encoding='cp1252')
```

```
In [3]: f_birth.head()
```

```
Out[3]:
```

	Date	Daily total female births in California, 1959
0	01-01-1959	35
1	02-01-1959	32
2	03-01-1959	30
3	04-01-1959	31
4	05-01-1959	44

```
In [4]: f_birth.tail()
```

```
Out[4]:
```

	Date	Daily total female births in California, 1959
360	27-12-1959	37
361	28-12-1959	52
362	29-12-1959	48
363	30-12-1959	55
364	31-12-1959	50

Data Preparation

```
In [5]: f_birth.isnull().sum()
```

```
Out[5]: Date                                0
Daily total female births in California, 1959  0
dtype: int64
```

```
In [6]: f_birth.describe()
```

```
Out[6]:
```

Daily total female births in California, 1959	
count	365.000000
mean	41.980822
std	7.348257
min	23.000000
25%	37.000000
50%	42.000000
75%	46.000000
max	73.000000

```
In [7]: f_birth.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 2 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Date                                365 non-null    object
1   Daily total female births in California, 1959  365 non-null    int64
dtypes: int64(1), object(1)
memory usage: 5.8+ KB
```

```
In [8]: f_birth = f_birth.groupby('Date')['Daily total female births in California, 1959']
```

```
In [9]: f_birth.head()
```

```
Out[9]:
```

	Date	Daily total female births in California, 1959
0	01-01-1959	35
1	01-02-1959	23
2	01-03-1959	35
3	01-04-1959	39
4	01-05-1959	32

```
In [10]: f_birth['Date'] = pd.to_datetime(f_birth['Date'],format='%d-%m-%Y')
```

```
In [11]: f_birth.rename(columns={'Daily total female births in California, 1959':'female_births'})
```

```
In [12]: f_birth.head()
```

```
Out[12]:
```

	Date	female_births
0	1959-01-01	35
1	1959-02-01	23
2	1959-03-01	35
3	1959-04-01	39
4	1959-05-01	32


```
In [13]: f_birth['female_births']
```

```
Out[13]:
```

0	35
1	23
2	35
3	39
4	32
	..
360	38
361	41
362	44
363	46
364	50


Name: female_births, Length: 365, dtype: int64

```
In [14]: f_birth['female_births'].plot()  
plt.show()
```




```
In [15]: a= f_birth.sample(20)
```

```
In [16]: sns.barplot(x='Date',y='female_births',data=a)
plt.xticks(rotation = 90)
plt.show()
```




```
In [17]: plt.figure(figsize=(10, 6))
sns.lineplot(x='Date', y='female_births', data=f_birth)
plt.xlabel('Date')
plt.ylabel('Female Births')
plt.xticks(rotation=90)
plt.show()
```



AutoCorrelation Plots

```
In [18]: from pandas.plotting import autocorrelation_plot as aplt
```

```
In [19]: aplt(f_birth['female_births'])  
plt.show()
```



stationarity and plot ACF and PACF

```
In [27]: # Function to test stationarity and plot ACF and PACF
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.ar_model import AutoReg
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
# Dickey-Fuller test
from statsmodels.tsa.stattools import adfuller
def test_stationarity(ts):

    result = adfuller(ts, autolag='AIC')
    print(f'ADF Statistic: {result[0]}')
    print(f'p-value: {result[1]}')
    print('Critical Values:')
    for key, value in result[4].items():
        print(f'    {key}: {value}')

# Plot ACF and PACF
plot_acf(ts, lags=20)
plot_pacf(ts, lags=20)
plt.show()
```



```
In [28]: # Test stationarity  
test_stationarity(f_birth['female_births'])
```

ADF Statistic: -16.680635544776198


p-value: 1.5174514677228693e-29

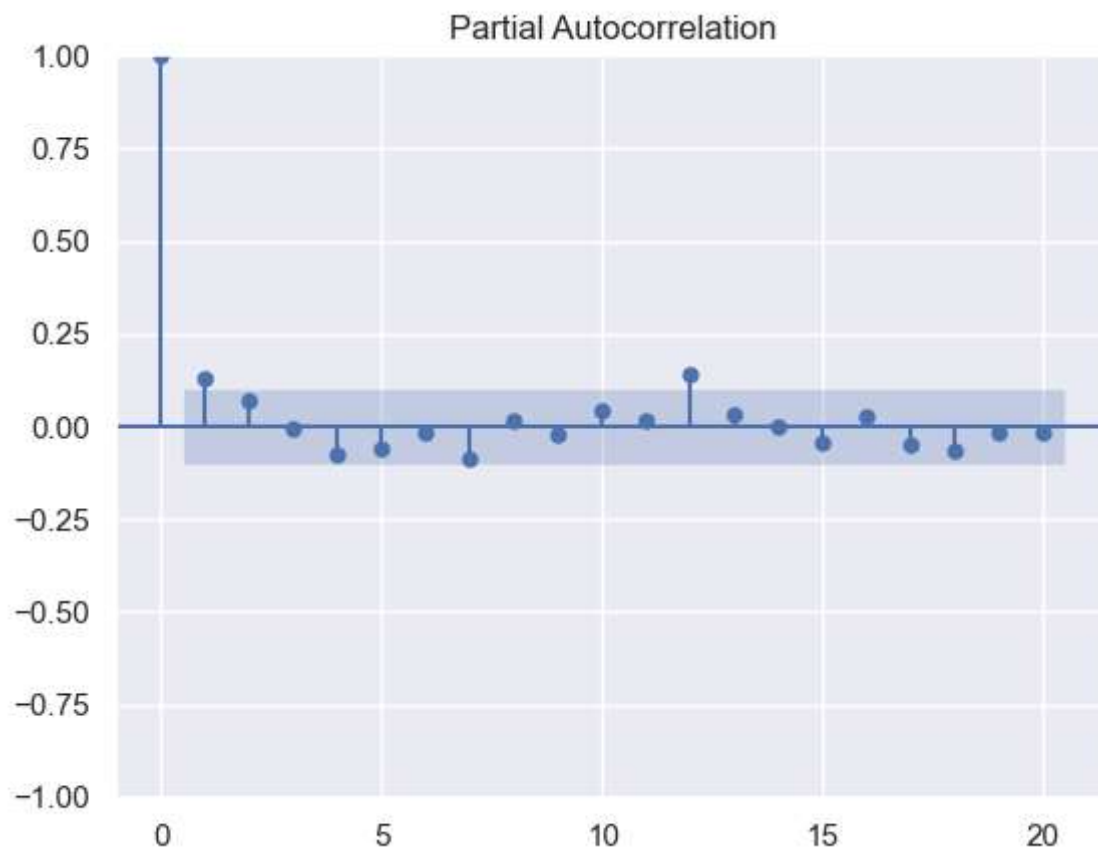
Critical Values:

1%: -3.4484434475193777

5%: -2.869513170510808

10%: -2.571017574266393





```
In [29]: # Split the data into training and testing sets
train_size = int(len(f_birth) * 0.8)
train, test = f_birth['female_births'][:train_size], f_birth['female_births'][t
```

```
In [30]: train_size
```

```
Out[30]: 292
```

```
In [31]: train
```

```
Out[31]: 0      35
1      23
2      35
3      39
4      32
..
287    38
288    37
289    45
290    34
291    37
Name: female_births, Length: 292, dtype: int64
```

```
In [32]: test
```

```
Out[32]: 292    37
          293    30
          294    36
          295    48
          296    44
          ..
          360    38
          361    41
          362    44
          363    46
          364    50
          Name: female_births, Length: 73, dtype: int64
```

```
In [33]: # AR Model
ar_model = AutoReg(train, lags=1).fit()
ar_pred = ar_model.predict(start=len(train), end=len(train) + len(test) - 1)
```

```
In [34]: # MA Model
ma_model = ARIMA(train, order=(0, 0, 1)).fit()
ma_pred = ma_model.predict(start=len(train), end=len(train) + len(test) - 1)
```

```
In [35]: # ARMA Model
arma_model = ARIMA(train, order=(1, 0, 1)).fit()
arma_pred = arma_model.predict(start=len(train), end=len(train) + len(test) -
```

```
In [36]: # ARIMA Model
arima_model = ARIMA(train, order=(1, 1, 1)).fit()
arima_pred = arima_model.predict(start=len(train), end=len(train) + len(test)
```

```
In [37]: # Forecasting with ARIMA
model = ARIMA(f_birth['female_births'], order=(1, 1, 1)).fit()
forecast_steps = 5 # Change this according to your needs
forecast = model.forecast(steps=forecast_steps, typ='levels')
```

```

In [38]: # Plotting results
plt.figure(figsize=(12, 8))

plt.subplot(2, 2, 1)
plt.plot(test, label='Actual')
plt.plot(ar_pred, label='AR Model')
plt.title('AR Model Forecast')
plt.legend()

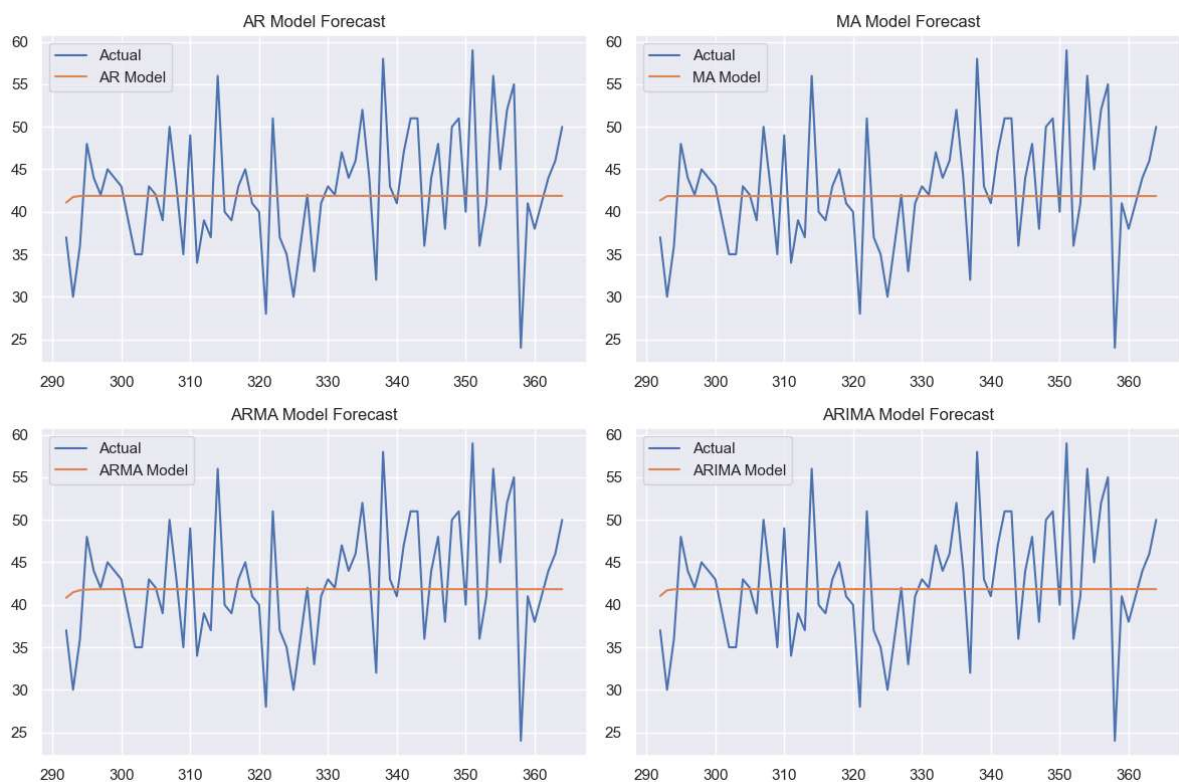
plt.subplot(2, 2, 2)
plt.plot(test, label='Actual')
plt.plot(ma_pred, label='MA Model')
plt.title('MA Model Forecast')
plt.legend()

plt.subplot(2, 2, 3)
plt.plot(test, label='Actual')
plt.plot(arma_pred, label='ARMA Model')
plt.title('ARMA Model Forecast')
plt.legend()


plt.subplot(2, 2, 4)
plt.plot(test, label='Actual')
plt.plot(arima_pred, label='ARIMA Model')
plt.title('ARIMA Model Forecast')
plt.legend()

plt.tight_layout()
plt.show()

```



```
In [39]: # Plot Forecasting with ARIMA
plt.figure(figsize=(10, 6))
plt.plot(f_birth['female_births'], label='Historical Data')
plt.plot(forecast, label='ARIMA Forecast', linestyle='dashed', color='red')
plt.title('ARIMA Forecasting')
plt.legend(loc="upper left")
plt.show()
```



Seasonal Decomposition

```
In [40]: from statsmodels.tsa.seasonal import seasonal_decompose

# Load specific forecasting tools
from statsmodels.tsa.statespace.sarimax import SARIMAX

from statsmodels.tsa.seasonal import seasonal_decompose, DecomposeResult
from pmdarima import auto_arima # for determining

# Load specific evaluation tools
from sklearn.metrics import mean_squared_error
from statsmodels.tools.eval_measures import rmse
```

```

In [41]: def plot_seasonal_decompose(result:DecomposeResult, dates:pd.Series=None, title:
x_values = dates if dates is not None else np.arange(len(result.observed))
cols = px.colors.qualitative.Pastel
return (
    make_subplots(
        rows=4,
        cols=1,
        subplot_titles=["Observed", "Trend", "Seasonal", "Residuals"],
    )
    .add_trace(
        go.Scatter(x=x_values, y=result.observed, mode="lines", name='Observed',
        row=1,
        col=1,
    )
    .add_trace(
        go.Scatter(x=x_values, y=result.trend, mode="lines", name='Trend',
        row=2,
        col=1,
    )
    .add_trace(
        go.Scatter(x=x_values, y=result.seasonal, mode="lines", name='Seasonal',
        row=3,
        col=1,
    )
    .add_trace(
        go.Scatter(x=x_values, y=result.resid, mode="lines", name='Residuals',
        row=4,
        col=1,
    )
    .update_layout(
        height=900, title=f'<b>{title}</b>', margin={'t':100}, title_x=0.5
    )
)

```

```
In [44]: seasonal_df = pd.DataFrame(f_birth.groupby(['Date'])['female_births'].apply("m
decomposition = seasonal_decompose(seasonal_df['female_births'], model='additi
fig = plot_seasonal_decompose(decomposition, dates=seasonal_df.index)

fig.update_layout(plot_bgcolor='white')
fig.update_yaxes(
    mirror=True,
    ticks='outside',
    showline=True,
    linecolor='black',
    gridcolor='lightgrey',
    title=''
)

fig.show()
```



```
In [45]: results = seasonal_decompose(f_birth["female_births"], model="multiplicative",
```

```
In [46]: fig, (ax1, ax2, ax3, ax4) = plt.subplots(4, 1, figsize=(12, 8))

# Plot the seasonal decomposition components on respective subplots
results.observed.plot(ax=ax1)
ax1.set_ylabel('Observed')

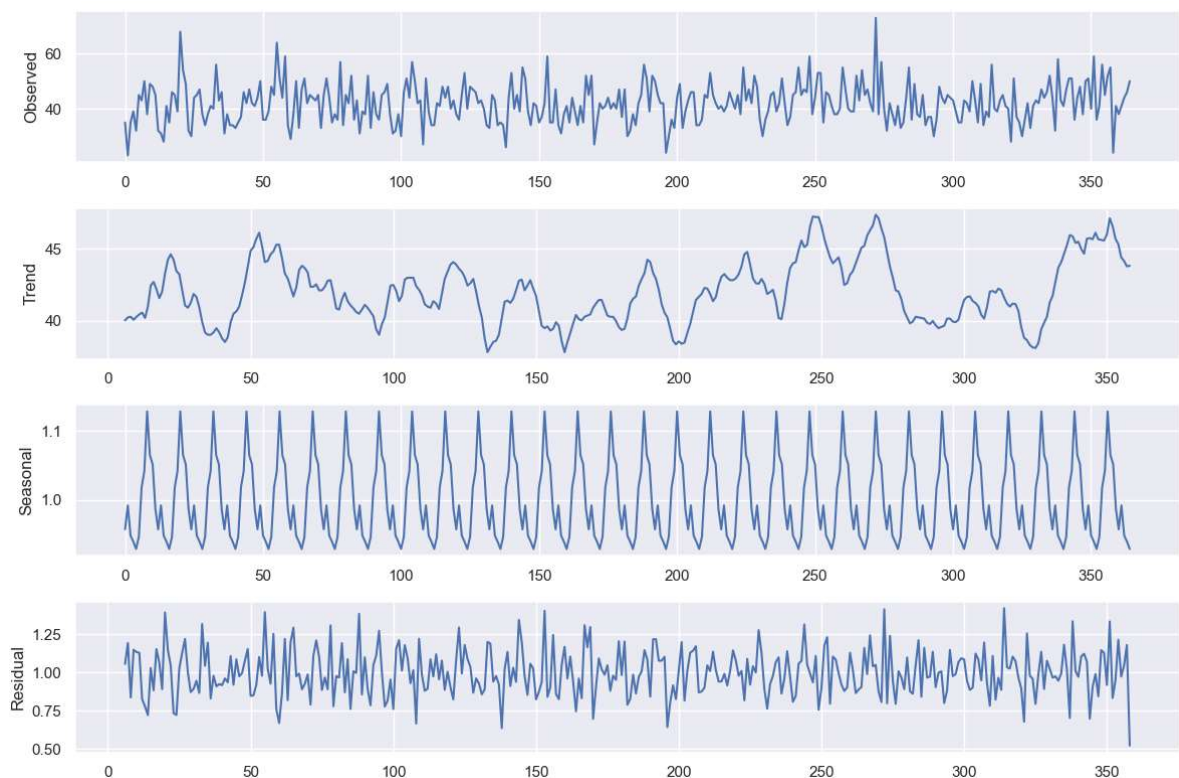
results.trend.plot(ax=ax2)
ax2.set_ylabel('Trend')

results.seasonal.plot(ax=ax3)
ax3.set_ylabel('Seasonal')

results.resid.plot(ax=ax4)
ax4.set_ylabel('Residual')

# Adjust layout for better presentation
fig.tight_layout()

# Display the plot
plt.show()
```



```
In [47]: #forecasting for Female_births
auto_arima(seasonal_df['female_births'], seasonal=True, m=12).summary()
```

Out[47]: SARIMAX Results

Dep. Variable:	y	No. Observations:	365
Model:	SARIMAX(1, 1, 1)x(1, 0, [], 12)	Log Likelihood	-1224.165
Date:	Fri, 22 Dec 2023	AIC	2456.330
Time:	15:16:57	BIC	2471.919
Sample:	01-01-1959 - 12-31-1959	HQIC	2462.526
Covariance Type:	opg		

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.1043	0.060	1.736	0.083	-0.013	0.222
ma.L1	-0.9544	0.019	-50.429	0.000	-0.991	-0.917
ar.S.L12	-0.1192	0.051	-2.327	0.020	-0.220	-0.019
sigma2	48.4909	3.247	14.935	0.000	42.127	54.855

Ljung-Box (L1) (Q):	0.02	Jarque-Bera (JB):	25.08
Prob(Q):	0.88	Prob(JB):	0.00
Heteroskedasticity (H):	0.94	Skew:	0.58
Prob(H) (two-sided):	0.72	Kurtosis:	3.54

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [48]: train = seasonal_df.iloc[:len(seasonal_df)-12]
test = seasonal_df.iloc[len(seasonal_df)-12:]
```

```
In [49]: model = SARIMAX(train['female_births'], order=(2, 1, 2), seasonal_order=(2, 0,
results = model.fit()
results.summary())
```

Out[49]: SARIMAX Results

Dep. Variable:	female_births	No. Observations:	353
Model:	SARIMAX(2, 1, 2)x(2, 0, [1], 12)	Log Likelihood	-1184.019
Date:	Fri, 22 Dec 2023	AIC	2384.039
Time:	15:17:00	BIC	2414.948
Sample:	01-01-1959 - 12-19-1959	HQIC	2396.339

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.2802	1.910	0.147	0.883	-3.464	4.024
ar.L2	0.0108	0.210	0.052	0.959	-0.401	0.423
ma.L1	-1.1350	1.912	-0.594	0.553	-4.883	2.613
ma.L2	0.1691	1.828	0.093	0.926	-3.413	3.752
ar.S.L12	0.5794	1.584	0.366	0.714	-2.524	3.683
ar.S.L24	0.0969	0.157	0.618	0.536	-0.210	0.404
ma.S.L12	-0.6950	1.591	-0.437	0.662	-3.813	2.423
sigma2	48.4428	3.342	14.493	0.000	41.892	54.994

Ljung-Box (L1) (Q): 0.02 **Jarque-Bera (JB):** 23.96
Prob(Q): 0.90 **Prob(JB):** 0.00
Heteroskedasticity (H): 0.96 **Skew:** 0.57
Prob(H) (two-sided): 0.84 **Kurtosis:** 3.58

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [50]: start = len(train)
end = len(train) + len(test) - 1

predictions = results.predict(start, end, typ='levels').rename('SARIMA Test Pr
```

```
In [51]: cols = px.colors.qualitative.Pastel
fig = go.Figure()
fig.add_trace(go.Scatter(y=test['female_births'],
                        mode='lines',
                        name='Real Values',
                        marker=dict(color=cols[0])
                        ))
fig.add_trace(go.Scatter(y=predictions,
                        mode='lines',
                        name='SARIMA Test Predictions',
                        marker=dict(color=cols[1])
                        ))
fig.update_layout(plot_bgcolor='white')
fig.update_yaxes(
    mirror=True,
    ticks='outside',
    showline=True,
    linecolor='black',
    gridcolor='lightgrey',
    title=''
)
fig.update_xaxes(
    dtick = 1
)
fig.show()
```

```
In [52]: error = rmse(test['female_births'], predictions)
print(f'RMSE Error: {error}')
```

RMSE Error: 6.874284881909711

```
In [53]: seasonal_df.describe().T[['mean', 'std']]
```

```
Out[53]:
```

	mean	std
female_births	41.980822	7.348257

```
In [54]: #Retraining the model on the entire dataset and forecast several years ahead
model = SARIMAX(seasonal_df['female_births'], order=(2, 1, 2), seasonal_order=
results = model.fit()
forecast = results.predict(start=len(seasonal_df), end=len(seasonal_df)+22, ty
```

```
In [55]: cols = px.colors.qualitative.Pastel
fig = go.Figure()
fig.add_trace(go.Scatter(x=seasonal_df.index, y=seasonal_df['female_births'],
                        mode='lines',
                        name='Real Values',
                        marker=dict(color=cols[0])
                        ))
fig.add_trace(go.Scatter(x=forecast.index, y=forecast,
                        mode='lines',
                        name='SARIMA Test Predictions',
                        marker=dict(color=cols[1])
                        ))
fig.update_layout(plot_bgcolor='white')
fig.update_yaxes(
    mirror=True,
    ticks='outside',
    showline=True,
    linecolor='black',
    gridcolor='lightgrey',
    title=''
)

fig.show()
```

Persistence Model

```
In [57]: f_birth['t']=f_birth['female_births'].shift(1)
```

```
In [58]: f_birth.head()
```

```
Out[58]:
```

	Date	female_births	t
0	1959-01-01	35	NaN
1	1959-02-01	23	35.0
2	1959-03-01	35	23.0
3	1959-04-01	39	35.0
4	1959-05-01	32	39.0

```
In [59]: train,test= f_birth[1:f_birth.shape[0]-7],f_birth[f_birth.shape[0]-7:]
```

```
In [60]: train.shape
```

```
Out[60]: (357, 3)
```

```
In [61]: test.shape
```


```
Out[61]: (7, 3)
```

```
In [62]: train.head()
```

```
Out[62]:
```

	Date	female_births	t
1	1959-02-01	23	35.0
2	1959-03-01	35	23.0
3	1959-04-01	39	35.0
4	1959-05-01	32	39.0
5	1959-06-01	45	32.0

```
In [63]: train['female_births'].plot(figsize=(16,6),fontsize=15)
plt.xlabel("Dates")
plt.ylabel('Female Birth')
plt.show()
```



Walk Forward Validation

```
In [64]: train_x,train_y = train['t'],train['female_births']
test_x,test_y = test['t'],test['female_births']
```

```
In [65]: predictions= test_x.copy()
```

```
In [66]: print(predictions)
print(test_y)
```


```
358    55.0
359    24.0
360    41.0
361    38.0
362    41.0
363    44.0
364    46.0
Name: t, dtype: float64
358     24
359     41
360     38
361     41
362     44
363     46
364     50
Name: female_births, dtype: int64
```



```
In [67]: from sklearn.metrics import mean_squared_error  
mse = mean_squared_error(test_y, predictions)  
mse
```

Out[67]: 185.28571428571428

```
In [68]: plt.plot(test_y)  
plt.plot(predictions, color='red')  
plt.show()
```



Random Forest Time Series Forecasting

```
In [69]: from sklearn.ensemble import RandomForestRegressor  
from sklearn.metrics import mean_absolute_error
```

```
In [70]: #transfer a timeseries data into a supervised Learning dataset
def series_to_supervised(data,n_in=1,n_out=1,dropnan=True):
    n_vars= 1 if type(data) is list else data.shape[1]
    df=pd.DataFrame(data)
    cols=list()
    #input sequence (t-n,...,t-1)
    for i in range(n_in,0,-1):
        cols.append(df.shift(i))
    #forecast sequence(t,t+1,...,t+n)

    for i in range(0,n_out):
        cols.append(df.shift(-i))
    #put it all together
    agg= pd.concat(cols,axis=1)
    #Drop rows with NAN values
    if dropnan:
        agg.dropna(inplace=True)
    return agg.values
```

```
In [71]: #Split univariate dataset into train/test sets -
def train_test_split(data, n_test):
    return data[:-n_test,:],data[-n_test:,:]
```

```
In [72]: #fit a Random Forest model and make a one step prediction
def random_forest_forecast(train, testX):
    #transform list into array
    train = np.asarray(train)
    #split into input and output columns
    trainX,trainY = train[:, :-1],train[:, -1]
    #fit model
    model = RandomForestRegressor(n_estimators=1000)
    model.fit(trainX,trainY)
    #make one step prediction
    y_pred = model.predict([testX])
    return y_pred[0]
```


```
In [73]: #Walk Forward validation for univariate data
def walk_forward_validation(data,n_test):
    predictions = list()
    #split dataset
    train, test = train_test_split(data,n_test)
    #seed history with training dataset
    history = [x for x in train]
    #step over each time step in the test set
    for i in range(len(test)):
        #split test row into input output columns
        testX, testY = test[i,:-1],test[i,-1]
        #fit model on history and make a prediction
        y_pred = random_forest_forecast(history, testX)
        #store forecast in list of predictions
        predictions.append(y_pred)
        #add actual observation to history for the next loop
        history.append(test[i])
        #summarize progress
        print(>expected=%.1f, predicted = %.1f'%(testY,y_pred))
    #estimate prediction error
    error = mean_absolute_error(test[:,-1],predictions)
    return error, test[:,-1],predictions
```

```
In [74]: #Load the dataset
series = pd.read_csv(r"E:\Data Science Projects\Project Files\Timeseries\Daily
                    header=0,index_col=0)
values = series.values
#transform the timeseries data set into supervised learning dataset
data= series_to_supervised(values,n_in=6)
```

```
In [75]: #evaluate
mae, y , y_pred = walk_forward_validation(data,12)
print("MAE:%3f"%mae)
```

```
>expected=42.0, predicted = 45.9
>expected=53.0, predicted = 43.6
>expected=39.0, predicted = 41.6
>expected=40.0, predicted = 38.2
>expected=38.0, predicted = 42.7
>expected=44.0, predicted = 47.5
>expected=34.0, predicted = 42.3
>expected=37.0, predicted = 36.6
>expected=52.0, predicted = 38.4
>expected=48.0, predicted = 41.6
>expected=55.0, predicted = 43.8
>expected=50.0, predicted = 45.3
MAE:5.873583
```

```
In [76]: #plot expected vs predicted
plt.plot(y,label='Expected')
plt.plot(y_pred,label = 'Predicted')
plt.legend()
plt.show()
```



```
In [ ]:
```